



GPU Workshop

Harvard



3 Ways to Accelerate Applications

Applications

Libraries

“Drop-in”
Acceleration

OpenACC
Directives

Easily Accelerate
Applications

Programming
Languages

Maximum
Flexibility


GPU Accelerated Libraries

“Drop-in” Acceleration for your Applications


Linear Algebra
FFT, BLAS,
SPARSE, Matrix




Numerical & Math
RAND, Statistics



Data Struct. & AI
Sort, Scan, Zero Sum



Visual Processing
Image & Video



3 Ways to Accelerate Applications

Applications

Libraries

“Drop-in”
Acceleration

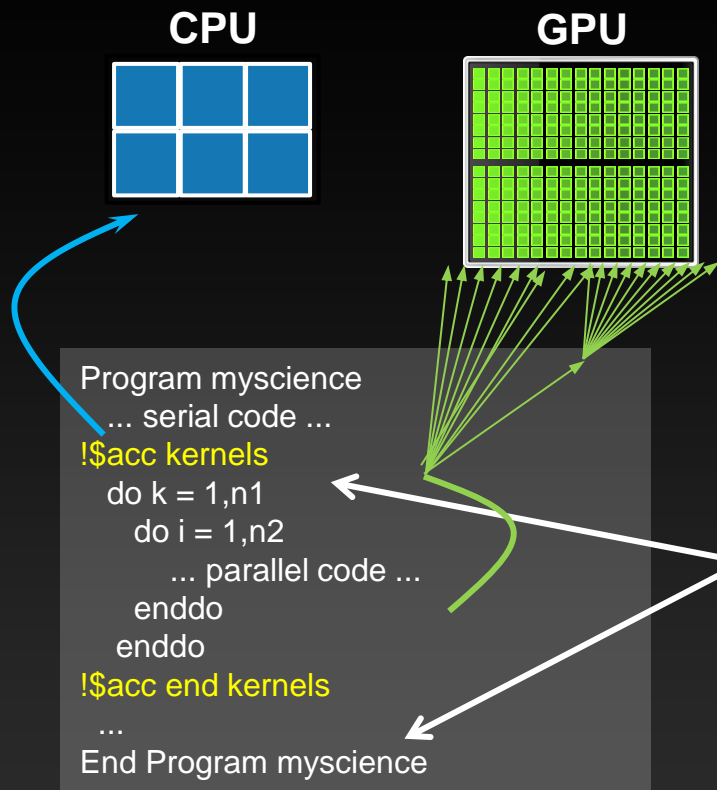
OpenACC
Directives

Easily Accelerate
Applications

Programming
Languages

Maximum
Flexibility

OpenACC Directives



OpenACC
Compiler
Hint

Simple Compiler hints

Compiler Parallelizes code

Works on many-core GPUs &
multicore CPUs

Your original
Fortran or C code

Familiar to OpenMP Programmers



OpenMP

CPU




```
main() {  
    double pi = 0.0; long i;  
  
    #pragma omp parallel for reduction(+:pi)  
    for (i=0; i<N; i++)  
    {  
        double t = (double)((i+0.05)/N);  
        pi += 4.0/(1.0+t*t);  
    }  
  
    printf("pi = %f\n", pi/N);  
}
```

OpenACC

CPU



GPU



```
main() {  
    double pi = 0.0; long i;  
  
    #pragma acc kernels  
    for (i=0; i<N; i++)  
    {  
        double t = (double)((i+0.05)/N);  
        pi += 4.0/(1.0+t*t);  
    }  
  
    printf("pi = %f\n", pi/N);  
}
```

Directives: Easy & Powerful



Real-Time Object Detection

Global Manufacturer of Navigation Systems



5x in 40 Hours

Valuation of Stock Portfolios using Monte Carlo


Global Technology Consulting Company



2x in 4 Hours

Interaction of Solvents and Biomolecules

University of Texas at San Antonio



5x in 8 Hours

“Optimizing code with directives is quite easy, especially compared to CPU threads or writing CUDA kernels. The most important thing is avoiding restructuring of existing code for production applications.”

-- Developer at the Global Manufacturer of Navigation Systems

A Very Simple Exercise: SAXPY



SAXPY in C

```
void saxpy(int n,
           float a,
           float *x,
           float *restrict y)
{
    #pragma acc kernels
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

...
// Perform SAXPY on 1M elements
saxpy(1<<20, 2.0, x, y);
...
```

SAXPY in Fortran


```
subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    !$acc kernels
    do i=1,n
        y(i) = a*x(i)+y(i)
    enddo
    !$acc end kernels
end subroutine saxpy

...
$ Perform SAXPY on 1M elements
call saxpy(2**20, 2.0, x_d, y_d)
...
```


GPU Architecture

GPU Architecture: Two Main Components

- **Global memory**
 - Analogous to RAM in a CPU server
 - Accessible by both GPU and CPU
 - Currently up to **12 GB**
 - **ECC on/off** option for Quadro and Tesla products
- **Streaming Multiprocessors (SMs)**
 - Perform the actual computations
 - Each SM has its own:
 - Control units, registers, execution pipelines, caches



GPU Architecture

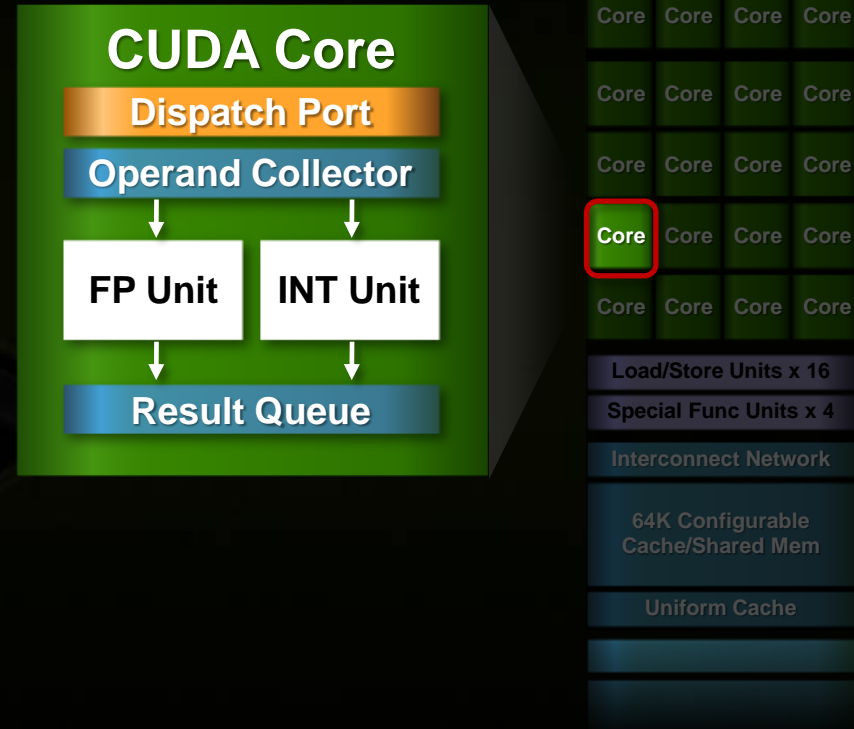
Streaming Multiprocessor (SM)

- Many CUDA Cores per SM
 - Architecture dependent
- Special-function units
 - cos/sin/tan, etc.
- Shared mem + L1 cache
- Thousands of 32-bit registers

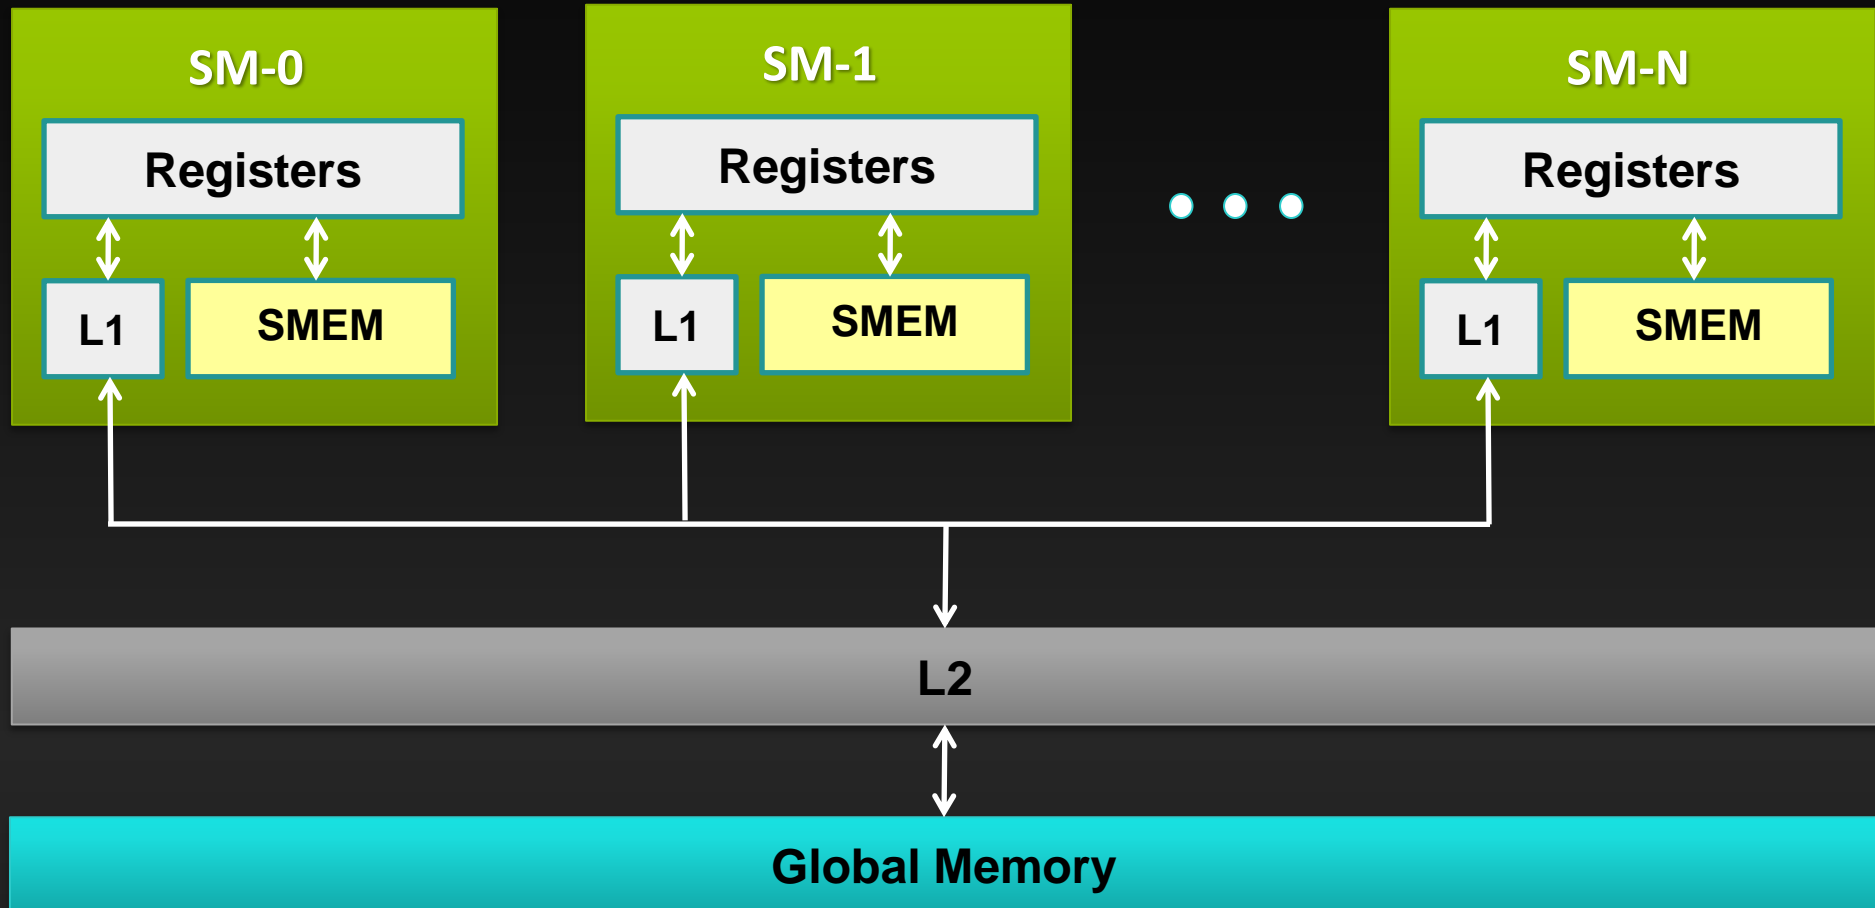


GPU Architecture: CUDA Core

- **Floating point & Integer unit**
 - **IEEE 754-2008 floating-point standard**
 - **Fused multiply-add (FMA) instruction for both single and double precision**
- **Logic unit**
- **Move, compare unit**
- **Branch unit**



GPU Memory Hierarchy Review



GPU Architecture:

Memory System on each SM

- **Extremely fast, but small, i.e., 10s of Kb**
- **Programmer chooses whether to use cache as L1 or Shared mem**
 - **L1**
 - Hardware-managed—used for things like register spilling.
 - Should NOT attempt to utilize like CPU caches.
 - High Aggregate bandwidth per GPU.
 - **Shared memory**
 - User-managed scratch-pad.
 - Useful if multiple threads need repeated access to the same data, or threads need to share data.
 - **Programmer MUST synchronize data accesses to avoid race conditions!**

GPU Architecture:

Memory System on each GPU board

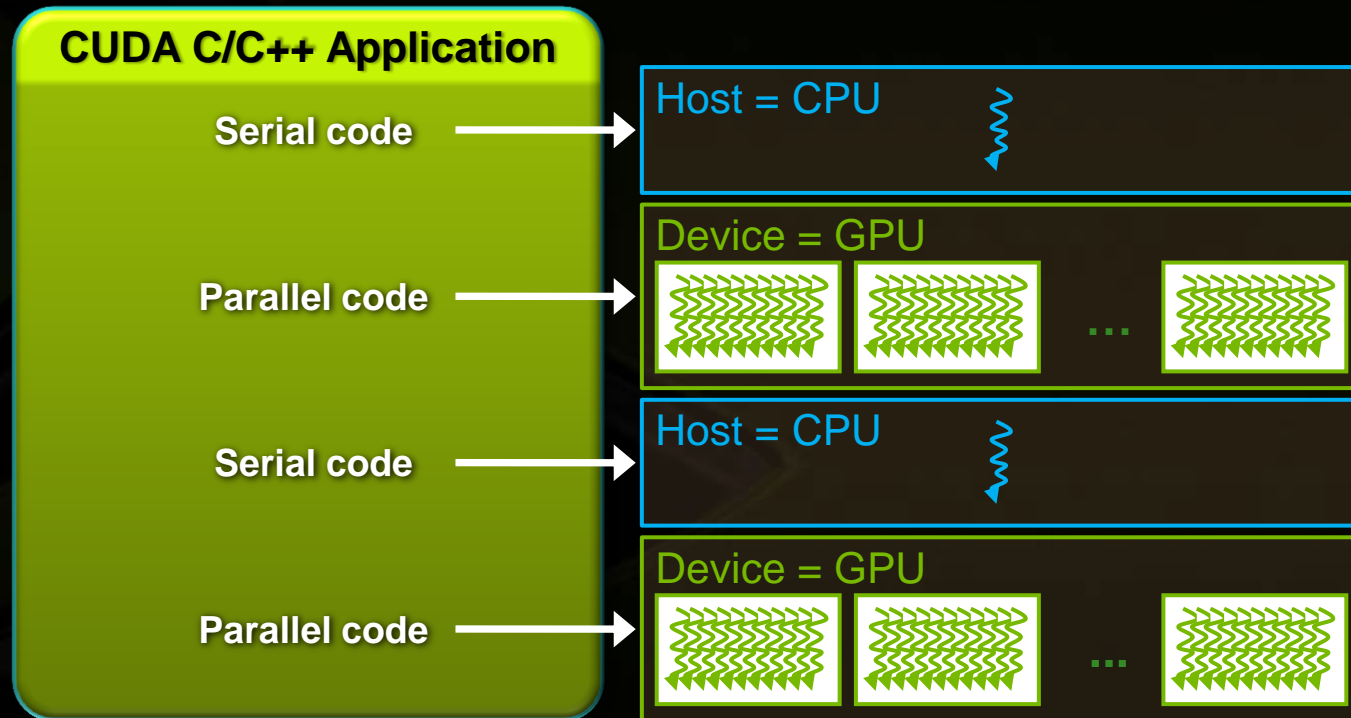
- **Unified L2 cache (100s of Kb)**
 - Fast, coherent data sharing across all cores in the GPU
- **ECC protection**
 - **DRAM**
 - ECC supported for GDDR5 memory
 - **All major internal memories are ECC protected**
 - Register file, L1 cache, L2 cache



CUDA PROGRAMMING MODEL

Anatomy of a CUDA C/C++ Application

- **Serial** code executes in a **Host** (CPU) thread
- **Parallel** code executes in many **Device** (GPU) threads across multiple processing elements



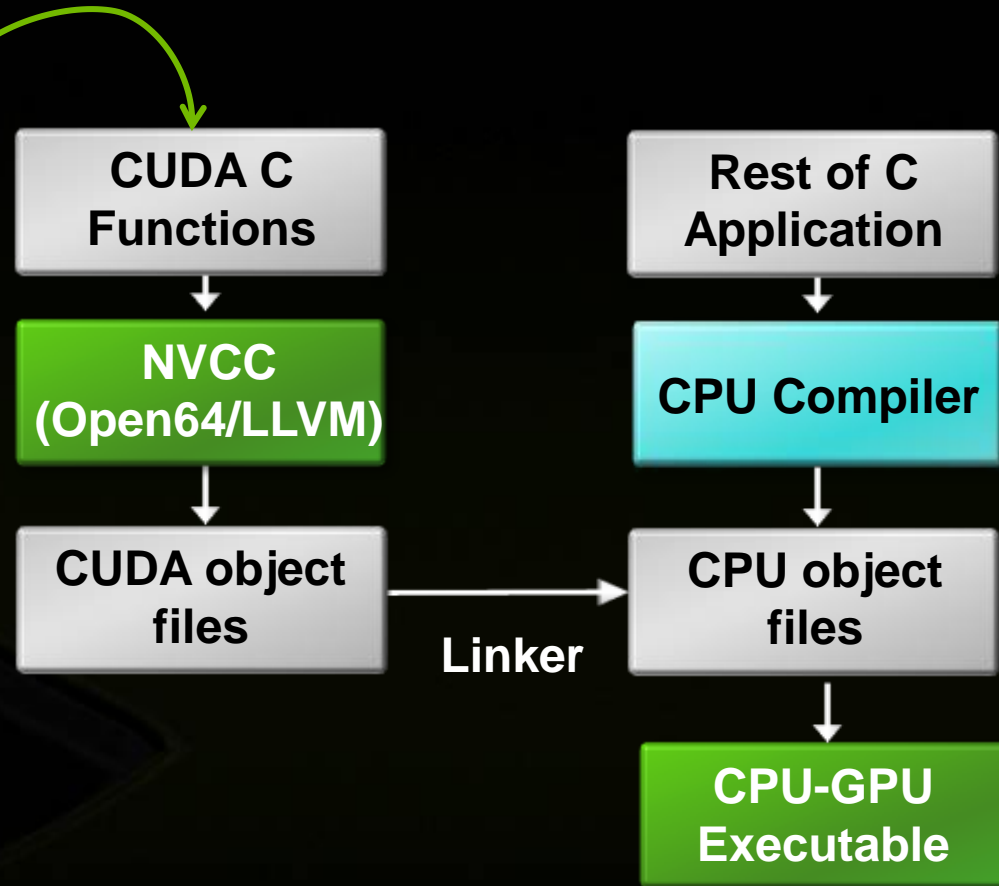
Compiling CUDA C Applications

```
void serial_function(... ) {  
    ...  
}  
void other_function(int ... ) {  
    ...  
}
```

```
void saxpy_serial(float ... ) {  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}
```

```
void main( ) {  
    float x;  
    saxpy_serial(..);  
    ...  
}
```

Modify into
Parallel
CUDA C code



CUDA C : C with a few keywords

```
void saxpy_serial(int n, float a, float *x, float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}
// Invoke serial SAXPY kernel
saxpy_serial(n, 2.0, x, y);
```

Standard C Code

```
__global__ void saxpy_parallel(int n, float a, float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}
// Invoke parallel SAXPY kernel with 256 threads/block
int nblocks = (n + 255) / 256;
saxpy_parallel<<<nblocks, 256>>>(n, 2.0, x, y);
```

Parallel C Code

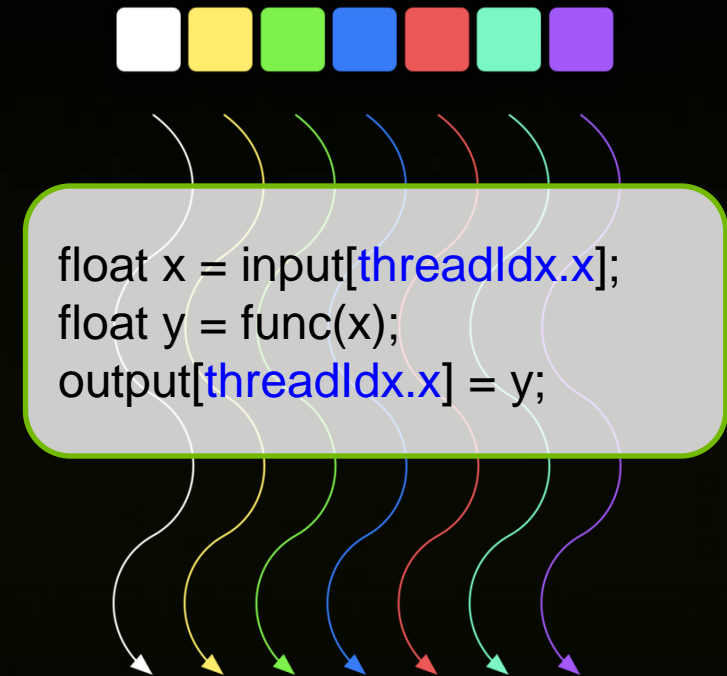
CUDA Kernels

- **Parallel portion of application: execute as a **kernel****
 - Entire GPU executes kernel, many threads
- **CUDA threads:**
 - Lightweight
 - Fast switching
 - 1000s execute simultaneously

CPU	Host	Executes functions
GPU	Device	Executes kernels

CUDA Kernels: Parallel Threads

- A **kernel** is a function executed on the GPU as an array of threads in parallel
- All threads execute the same code, can take different paths
- Each thread has an ID
 - Select input/output data
 - Control decisions



CUDA Kernels: Subdivide into Blocks




CUDA Kernels: Subdivide into Blocks




- Threads are grouped into **blocks**

CUDA Kernels: Subdivide into Blocks



- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**

CUDA Kernels: Subdivide into Blocks



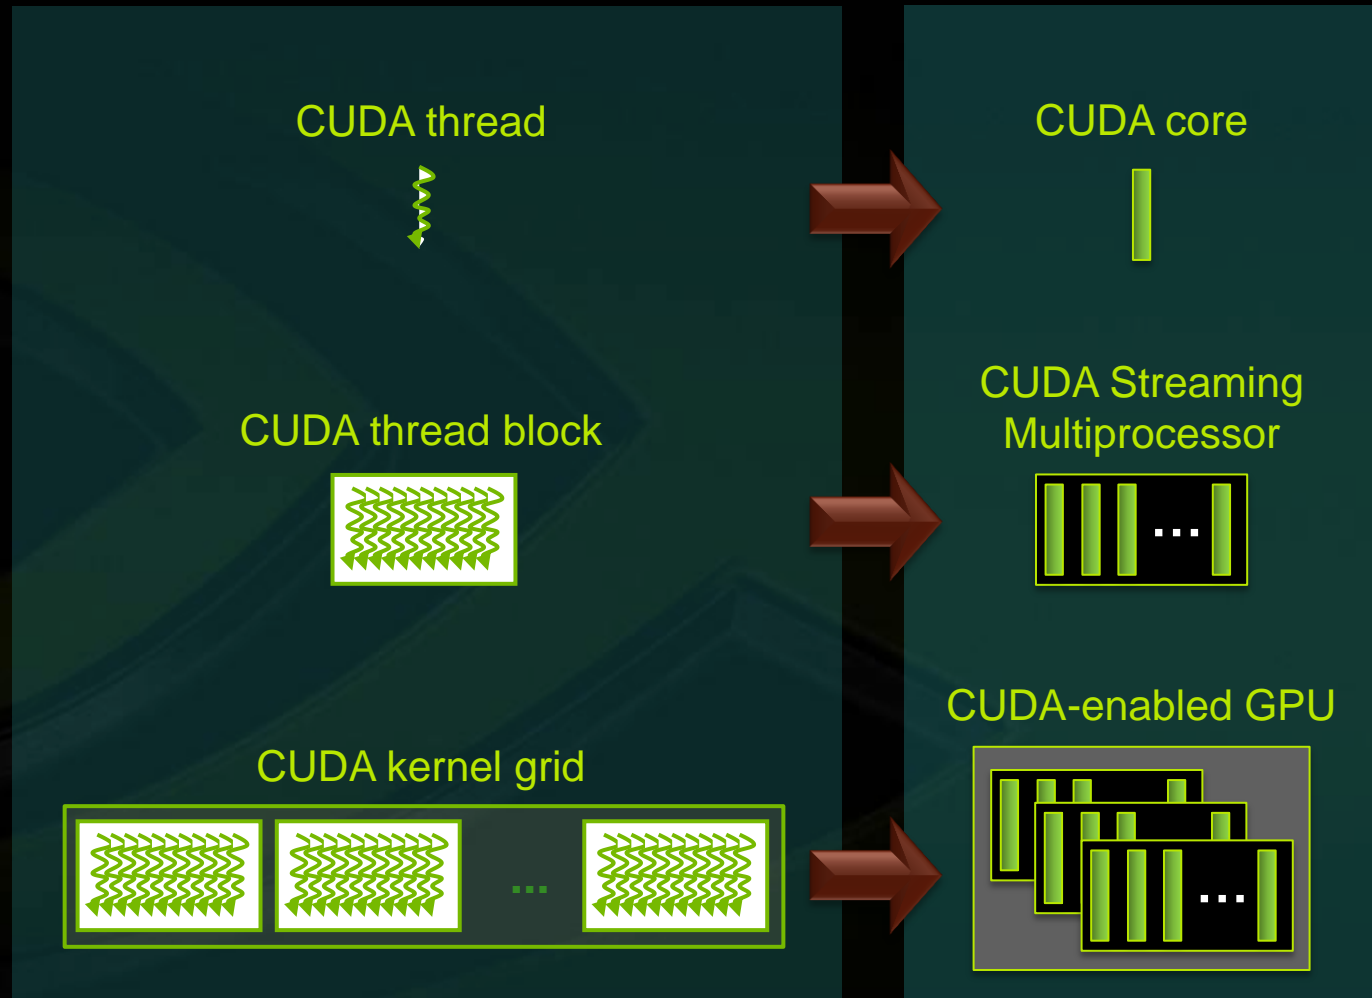
- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**
- A **kernel** is executed as a **grid of blocks of threads**

CUDA Kernels: Subdivide into Blocks



- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**
- A **kernel** is executed as a **grid** of **blocks** of **threads**


Kernel Execution



- Each thread is executed by a core
- Each block is executed by one SM and does not migrate
- Several concurrent blocks can reside on one SM depending on the blocks' memory requirements and the SM's memory resources
- Each kernel is executed on one device
- Multiple kernels can execute on a device at one time

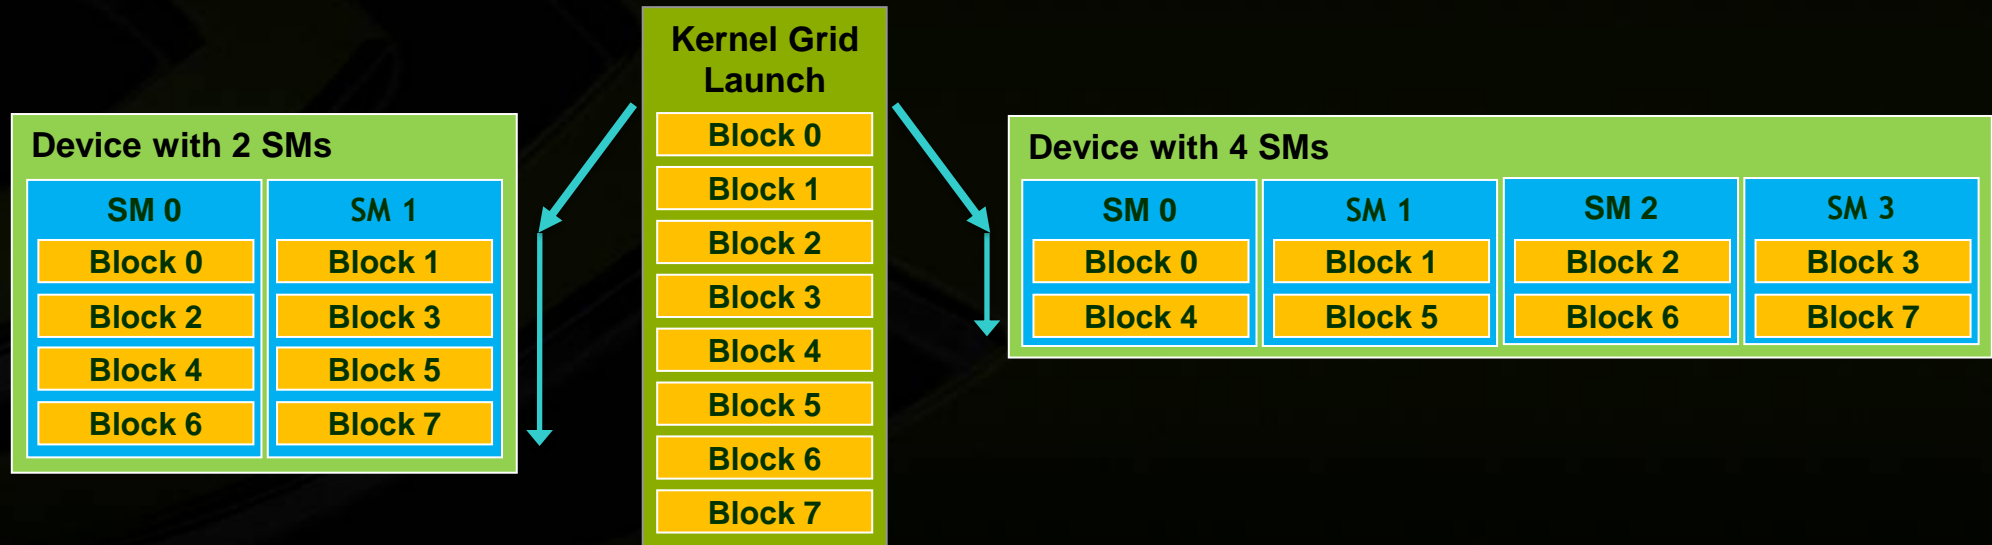
Thread blocks allow cooperation

- Threads may need to cooperate:
 - Cooperatively load/store blocks of memory all will use
 - Share results with each other or cooperate to produce a single result
 - Synchronize with each other



Thread blocks allow scalability


- Blocks can execute in any order, concurrently or sequentially
- This independence between blocks gives scalability:
 - A kernel scales across any number of SMs



MEMORY SYSTEM HIERARCHY


Memory hierarchy

- Thread:
 - Registers




Memory hierarchy

- Thread:
 - Registers
 - Local memory



Memory hierarchy


- **Thread:**
 - **Registers**
 - **Local** memory
- **Block of threads:**
 - **Shared** memory



Memory hierarchy : Shared memory


```
__shared__ int a[SIZE];
```

- Allocated per thread block, same lifetime as the block
- Accessible by any thread in the block
- Several uses:
 - Sharing data among threads in a block
 - User-managed cache (reducing gmem accesses)




Memory hierarchy

- Thread:
 - Registers
 - Local memory
- Block of threads:
 - Shared memory
- All blocks:
 - Global memory



Memory hierarchy : Global memory

- Accessible by all threads of any kernel
- Data lifetime: from allocation to deallocation by host code
 - `cudaMalloc (void ** pointer, size_t nbytes)`
 - `cudaMemset (void * pointer, int value, size_t count)`
 - `cudaFree (void* pointer)`



CUDA memory management



Memory Spaces

CPU and GPU have separate memory spaces

- Data is moved across PCIe bus
- Use functions to allocate/set/copy memory on GPU
 - Very similar to corresponding C functions

Pointers are just addresses

- Can't tell from the pointer value whether the address is on CPU or GPU
 - Must use **cudaPointerGetAttributes(...)**
- Must exercise care when dereferencing:
 - Dereferencing CPU pointer on GPU will likely crash
 - Dereferencing GPU pointer on CPU will likely crash

GPU Memory Allocation / Release

Host (CPU) manages device (GPU) memory

- `cudaMalloc (void ** pointer, size_t nbytes)`
- `cudaMemset (void * pointer, int value, size_t count)`
- `cudaFree (void* pointer)`

```
int n = 1024;  
int nbytes = 1024*sizeof(int);  
int * d_a = 0;  
cudaMalloc( (void**)&d_a, nbytes );  
cudaMemset( d_a, 0, nbytes);  
cudaFree(d_a);
```

Note: Device memory from GPU point of view is also referred to as global memory.

Data Copies

```
cudaMemcpy( void *dst, void *src, size_t nbytes,  
            enum cudaMemcpyKind direction);
```

- returns after the copy is complete
- blocks CPU thread until all bytes have been copied
- doesn't start copying until previous CUDA calls complete

```
enum cudaMemcpyKind
```

- cudaMemcpyHostToDevice
- cudaMemcpyDeviceToHost
- cudaMemcpyDeviceToDevice

Non-blocking memcpyes are provided



BASIC KERNELS AND EXECUTION

CUDA Programming Model revisited

- Parallel code (kernel) is launched and executed on a device by many threads
- Threads are grouped into thread blocks
- Parallel code is written for a thread
 - Each thread is free to execute a unique code path
 - Built-in thread and block ID variables

Thread Hierarchy

- Threads launched for a parallel section are partitioned into thread blocks
 - Grid = all blocks for a given launch
- Thread block is a group of threads that can:
 - Synchronize their execution
 - Communicate via shared memory

IDs and Dimensions

Threads

- 3D IDs, unique within a block

Blocks


- 2D IDs, unique within a grid

Dimensions set at launch time

- Can be unique for each grid

Built-in variables

- threadIdx, blockIdx
- blockDim, gridDim



(Continued)

IDs and Dimensions

Threads

- 3D IDs, unique within a block

Blocks

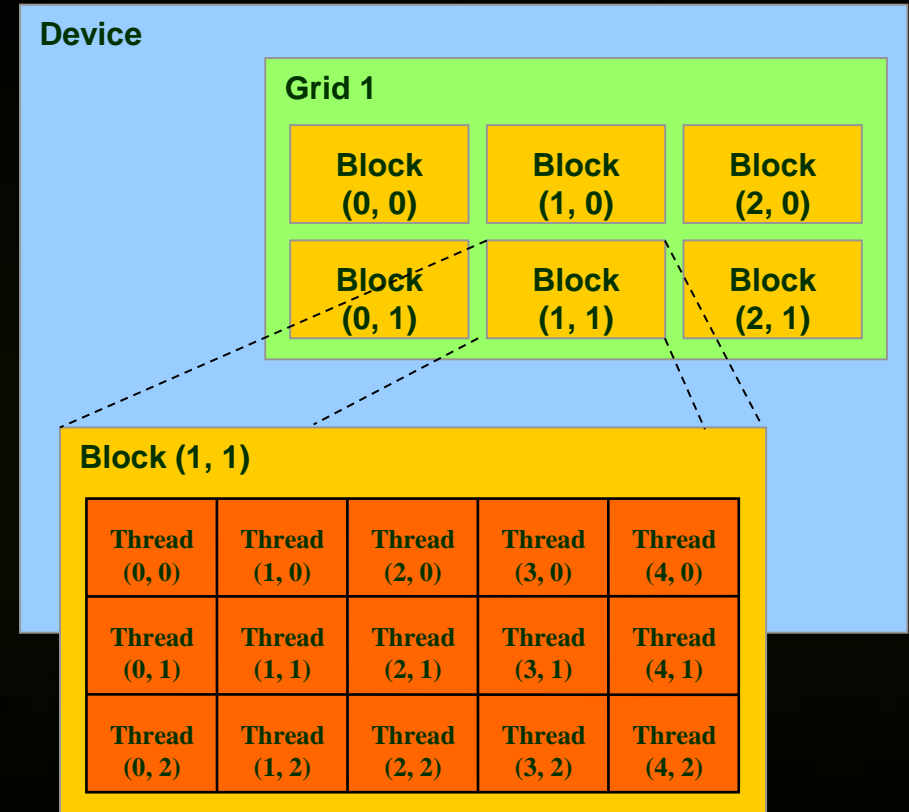
- 2D IDs, unique within a grid

Dimensions set at launch time

- Can be unique for each grid

Built-in variables

- threadIdx, blockIdx
- blockDim, gridDim



Launching Kernels on GPU

Launch parameters (triple chevron <<<>>> notation)

- grid dimensions (up to 2D), dim3 type
- thread-block dimensions (up to 3D), dim3 type
- shared memory: number of bytes per block
 - for extern smem variables declared without size
 - Optional, 0 by default
- stream ID
 - Optional, 0 by default

```
dim3 grid(16, 16);  
dim3 block(16,16);  
kernel<<<grid, block, 0, 0>>>(...);  
kernel<<<32, 512>>>(...);
```

GPU kernel execution

- Kernel launches on a grid of blocks, <<<grid,block>>>(arg1,...)
- Each block is launched on one SM
 - A block is divided into warps of 32 threads each (think 32-way vector)
 - Warps in a block are scheduled and executed.
 - All threads in a warp execute same instruction simultaneously (think SIMD)
 - Number of blocks/SM determined by resources required by the block
 - Registers, shared memory, total warps, etc.
- Block runs to completion on SM it started on, no migration.

Warps (The rest of the story...)



A thread block consists of 32-thread warps


A warp is executed physically in parallel (SIMD) on a multiprocessor

Execution Model


Software



Thread



Thread Block




Grid

Hardware



Scalar
Processor



Multiprocessor



Device

Threads are executed by scalar processors

Thread blocks are executed on multiprocessors

Thread blocks do not migrate

Several concurrent thread blocks can reside on one multiprocessor - limited by multiprocessor resources (shared memory and register file)

A kernel is launched as a grid of thread blocks

Blocks Must Be Independent

Any possible interleaving of blocks should be valid

- presumed to run to completion without pre-emption
- can run in any order
- can run concurrently OR sequentially

Blocks may coordinate but not synchronize

- shared queue pointer: OK
- shared lock: BAD ... any dependence on order easily deadlocks

Independence requirement gives scalability

Blocks Must Be Independent

- Facilitates scaling of the same code across many devices

